

Financial Statistics - Coursework

Group 30

2025-11-23

Table of contents

1	Question 1 - CAPM	3
1.1	1. Prepare the Dataset	3
1.2	2. Estimate linear regression model	4
1.3	3. F-test for $H_0: \beta_1 = \beta_2$	5
1.4	4. t-test for $H_0: \beta_1 = 0$	5
2	Question 2 - probability of a positive asset return	6
3	Question 3 - CIR model for the term structure of interest rate	10

Table of contents

List of Figures

List of Tables

```
# Load libraries
library(readxl)
library(lubridate)
library(readr)
```

1 Question 1 - CAPM

1.1 1. Prepare the Dataset

```
# Use the dataset data coursework Q1
data_q1 <- read_excel("data_coursework_Q1.xls")

data_q1$date <- make_date(year = data_q1$year_, month = data_q1$month_, day = 1)
data_q1$IBM <- parse_number(data_q1$IBM)
df_q1 <- data_q1[, c("date", "SP500", "IBM", "1-month Tbill")]

colnames(df_q1)[colnames(df_q1) == "1-month Tbill"] <- "Rf"
colnames(df_q1)[colnames(df_q1) == "IBM"] <- "Ri"

df_q1$Rm <- (df_q1$SP500 / dplyr::lag(df_q1$SP500)) - 1
df_q1$Rf <- df_q1$Rf / 12

head(df_q1, 30)
```

A tibble: 30 × 5

date <date>	SP500 <dbl>	Ri <dbl>	Rf <dbl>	Rm <dbl>
1960-01-01	58.03	NA	0.02750000	NA
1960-02-01	55.78	NA	0.02416667	-0.038773048
1960-03-01	55.02	NA	0.02916667	-0.013624955
1960-04-01	55.73	NA	0.01583333	0.012904398
1960-05-01	55.22	NA	0.02250000	-0.009151265
1960-06-01	57.26	NA	0.02000000	0.036943137
1960-07-01	55.84	NA	0.01083333	-0.024799162
1960-08-01	56.51	NA	0.01416667	0.011998567
1960-09-01	54.81	NA	0.01333333	-0.030083171
1960-10-01	53.73	NA	0.01833333	-0.019704433
1960-11-01	55.47	NA	0.01083333	0.032384143
1960-12-01	56.80	NA	0.01333333	0.023976924
1961-01-01	59.72	NA	0.01583333	0.051408451
1961-02-01	62.17	NA	0.01166667	0.041024782
1961-03-01	64.12	NA	0.01666667	0.031365610
1961-04-01	65.83	NA	0.01416667	0.026668746
1961-05-01	66.50	NA	0.01500000	0.010177731
1961-06-01	65.62	NA	0.01666667	-0.013233083
1961-07-01	65.44	NA	0.01500000	-0.002743066
1961-08-01	67.79	NA	0.01166667	0.035910758
1961-09-01	67.26	NA	0.01416667	-0.007818262
1961-10-01	68.00	NA	0.01583333	0.011002081
1961-11-01	71.08	NA	0.01250000	0.045294118
1961-12-01	71.74	NA	0.01583333	0.009285312
1962-01-01	69.07	2.69	0.02000000	-0.037217731
1962-02-01	70.22	2.66	0.01666667	0.016649776
1962-03-01	70.29	2.64	0.01666667	0.000996867
1962-04-01	68.05	2.25	0.01833333	-0.031867976
1962-05-01	62.99	1.95	0.02000000	-0.074357090
1962-06-01	55.63	1.68	0.01666667	-0.116843943

```

# Construct the measures

# excess return for IBM (Ri - Rf)
df_q1$ri_rf <- df_q1$Ri - df_q1$Rf
# excess market return (Rm - Rf)
df_q1$rm_rf <- df_q1$Rm - df_q1$Rf
# squared market excess return
df_q1$rm_rf_sq <- df_q1$rm_rf^2

# D_t = 1 if market excess return > 0, 0 otherwise
df_q1$D <- ifelse(df_q1$rm_rf > 0, 1, 0)
df_q1$D_rm <- df_q1$D * df_q1$rm_rf
df_q1$Dinv_rm <- (1 - df_q1$D) * df_q1$rm_rf

```

1.2 2. Estimate linear regression model

```

# Model 1: Standard CAPM
# ri_rf = alpha + beta * rm_rf + u
capm_model <- lm(ri_rf ~ rm_rf, data = df_q1)
# show summary
summary(capm_model)

# - Model 2: Extended CAPM
# ri_rf = alpha + beta1*D_rm + beta2*Dinv_rm + beta3*rm_rf_sq + u
ext_model <- lm(ri_rf ~ D_rm + Dinv_rm + rm_rf_sq, data = df_q1)
# show summary
summary(ext_model)

```

Call:

```
lm(formula = ri_rf ~ rm_rf, data = df_q1)
```

Residuals:

Min	1Q	Median	3Q	Max
-8.878	-4.733	-2.628	7.244	19.250

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	9.8675	0.5256	18.775	<2e-16 ***
rm_rf	-5.0450	9.0589	-0.557	0.578

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.112 on 346 degrees of freedom

(24 observations deleted due to missingness)

Multiple R-squared: 0.0008956, Adjusted R-squared: -0.001992

F-statistic: 0.3101 on 1 and 346 DF, p-value: 0.578

Call:

```
lm(formula = ri_rf ~ D_rm + Dinv_rm + rm_rf_sq, data = df_q1)
```

Residuals:

Min	1Q	Median	3Q	Max
-9.040	-4.880	-2.794	6.682	19.697

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	9.0742	0.8315	10.913	<2e-16 ***
D_rm	18.0443	51.6197	0.350	0.7269

```
Dinv_rm      -49.9093    28.6993   -1.739    0.0829 .
rm_rf_sq     -350.2208   196.2948   -1.784    0.0753 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 7.097 on 344 degrees of freedom
 (24 observations deleted due to missingness)
 Multiple R-squared: 0.01081, Adjusted R-squared: 0.002179
 F-statistic: 1.253 on 3 and 344 DF, p-value: 0.2906

1.3 3. F-test for $H_0: \beta_1 = \beta_2$

```
# Restricted model: force beta1 = beta2
# ri_rf = alpha + beta * rm_rf + beta3*rm_rf_sq + u
restricted_model <- lm(ri_rf ~ rm_rf + rm_rf_sq, data = df_q1)

# F-test between restricted and unrestricted models
anova(restricted_model, ext_model)
```

A anova: 2×6

	Res.Df <dbl>	RSS <dbl>	Df <dbl>	Sum of Sq <dbl>	F <dbl>	Pr(>F) <dbl>
1	345	17373.33	NA	NA	NA	NA
2	344	17328.28	1	45.04091	0.8941493	0.3450191

1.3.1 Summary

The F-test compares the restricted symmetric CAPM model with the unrestricted asymmetric beta specification. The test statistic is $F = 0.894$ with a corresponding p-value of 0.345. Because the p-value is well above conventional significance levels (1%, 5%, or 10%), we fail to reject the null hypothesis that $\beta_1 = \beta_2$.

This implies that the market beta of IBM does not differ significantly between up-market periods ($rm_rf > 0$) and down-market periods ($rm_rf \leq 0$). In other words, the data do not provide evidence of return asymmetry or different systematic risk exposures across market states. The standard CAPM with a single beta coefficient appears sufficient for this sample.

1.4 4. t-test for $H_0: \alpha = 0$

```
# extract summary table
capm_summary <- summary(capm_model)$coefficients

# alpha estimate, t-value, p-value
alpha_est <- capm_summary["(Intercept)", "Estimate"]
alpha_t <- capm_summary["(Intercept)", "t value"]
alpha_p <- capm_summary["(Intercept)", "Pr(>|t|)"]

alpha_est
alpha_t
alpha_p
```

9.86746894909065

18.7745349820389

1.00130758429436e-54

1.4.1 Summary

The CAPM regression shows that the estimated alpha (intercept) for IBM is 9.87, with a t-statistic of 18.77 and a p-value effectively 0. This indicates that the abnormal return (Jensen's alpha) is highly significant, meaning

IBM delivered a large positive excess return that cannot be explained by market movements alone during the sample period. The null hypothesis $= 0$ is strongly rejected.

2 Question 2 - probability of a positive asset return

```
# Load the required package
library(readxl)
library(utils)

# Sub-question 2a: OLS regression over the whole sample period
# Load data
data <- read_excel("data_coursework_Q2.xls", sheet = 1)
data <- na.omit(data)
head(data)

# Separate response variable and regressors
returns <- data$ESR
regressors <- data[, -which(names(data) == "ESR" | names(data) == "mth")]

# Create binary response variable
Y <- ifelse(returns > 0, 1, 0)

# Prepare lead-lag data
Y_lead <- Y[-1]
X_lag <- regressors[-length(Y), ]

reg_data <- data.frame(Y_lead, X_lag)

ols_model <- lm(Y_lead ~ ., data = reg_data)
summary(ols_model)
```

A tibble: 6 × 12

mth <chr>	ESR <dbl>	USR <dbl>	EIF <dbl>	ERECB <dbl>	EFX <dbl>	EDY <dbl>	ESP <dbl>	UIF <dbl>	URR <dbl>	UDY <dbl>	UOIL <dbl>
31- gen-85	0.0559	0.0690	0.0209	1.6094	0.0025	1.1663	2.206	0.0396	2.1668	1.4951	0.0210
31106	0.0331	0.0066	0.0235	1.6582	0.0556	1.1474	2.374	0.0347	2.1679	1.4929	0.0094
31137	-	-	0.0249	1.6582	-	1.1474	2.142	0.0354	2.1494	1.5041	0.0582
	0.0027	0.0089			0.0813						
31167	0.0226	-	0.0250	1.6582	0.0049	1.1569	2.023	0.0372	2.1494	1.5173	-
		0.0057									0.0269
31- mag- 85	0.0424	0.0504	0.0236	1.6582	-	1.1282	1.856	0.0352	2.0334	1.4679	0.0108
					0.0127						
30- giu-85	0.0263	0.0116	0.0193	1.6582	-	1.1019	1.753	0.0351	2.0732	1.4563	-
					0.0099						0.0347

Call:

```
lm(formula = Y_lead ~ ., data = reg_data)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-0.9505 -0.4431  0.1862  0.3995  0.9720
```

Coefficients:

```
      Estimate Std. Error t value Pr(>|t|)
```

(Intercept)	0.44284	0.30703	1.442	0.150325	
USR	1.04566	0.62401	1.676	0.094915	.
EIF	8.59691	3.88055	2.215	0.027541	*
ERECB	-0.62801	0.15006	-4.185	3.83e-05	***
EFX	0.94970	0.88001	1.079	0.281431	
EDY	0.58948	0.28929	2.038	0.042525	*
ESP	-0.06447	0.03875	-1.664	0.097318	.
UIF	-12.32022	3.70983	-3.321	0.001017	**
URR	0.42052	0.09377	4.485	1.07e-05	***
UDY	0.10158	0.16057	0.633	0.527482	
UOIL	-0.95801	0.28501	-3.361	0.000884	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4675 on 279 degrees of freedom
Multiple R-squared: 0.1352, Adjusted R-squared: 0.1042
F-statistic: 4.363 on 10 and 279 DF, p-value: 1.081e-05

The model is statistically significant ($p < 0.001$), and although the adjusted R^2 of around 10% indicates a modest explanatory power, it suggests that the included predictors have a meaningful impact on the likelihood of achieving a positive return. Several predictors—including EIF, EDY, and URR—significantly increase the probability of a positive return, while variables such as ERECB, UIF, and UOIL significantly reduce it.

```
# sub-question 2b: Compute Z(alpha=1/2) for all possible
# models and search for best predictors combination
# Compute predicted probabilities
P_hat <- predict(ols_model, type = "response")

# Construct Z_t(alpha=1/2)
alpha <- 0.5
returns_t <- returns[-1]
Z_t <- ifelse((P_hat > alpha & returns_t > 0) |
             (P_hat <= alpha & returns_t <= 0), 1, 0)

# overall hit ratio
Z_alpha <- mean(Z_t)
cat("Z(alpha = 0.5) for full model =", Z_alpha, "\n")

# Search for best combination of predictors that maximizes Z(1/2)
predictor_names <- names(X_lag)
k <- length(predictor_names)

cat("Number of predictors:", k, "\n")

compute_Z <- function(vars) {
  if (length(vars) == 0) {
    model <- lm(Y_lead ~ 1, data = reg_data)
  } else {
    fmla <- as.formula(paste("Y_lead ~", paste(vars, collapse = "+")))
    model <- lm(fmla, data = reg_data)
  }
  Ph <- predict(model, type = "response")
  Zt <- ifelse((Ph > alpha & returns_t > 0) |
             (Ph <= alpha & returns_t <= 0), 1, 0)
  return(list(Z = mean(Zt), model = model, vars = vars))
}

best <- list(Z = -Inf, vars = NULL, model = NULL)

for (m in 1:k) {
```

```

combs <- combn(predictor_names, m, simplify = FALSE)
for (vars in combs) {
  res <- compute_Z(vars)
  if (res$Z > best$Z) {
    best <- res
  }
}

# compare with intercept-only model
res0 <- compute_Z(character(0))
if (res0$Z > best$Z) {
  best <- res0
}
cat("Best Z(alpha=0.5) =", best$Z, "\n")
cat("Best predictor combination:\n")
print(best$vars)

```

```

Z(alpha = 0.5) for full model = 0.6655172
Number of predictors: 10
Best Z(alpha=0.5) = 0.6896552
Best predictor combination:
[1] "USR"  "EIF"  "ERECB" "EFX"  "EDY"  "ESP"  "UIF"  "URR"

print(summary(best$model))

```

```

Call:
lm(formula = fmla, data = reg_data)

```

```

Residuals:
    Min       1Q   Median       3Q      Max
-0.9325 -0.4699  0.2351  0.4077  0.7228

```

```

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.29068    0.19988   1.454 0.146976
USR          1.23219    0.63004   1.956 0.051488 .
EIF          5.83448    3.58057   1.629 0.104331
ERECB       -0.53419    0.12939  -4.129 4.82e-05 ***
EFX          1.00195    0.89248   1.123 0.262544
EDY          0.70824    0.18753   3.777 0.000194 ***
ESP         -0.06407    0.03804  -1.684 0.093198 .
UIF         -9.83058    3.69086  -2.663 0.008180 **
URR          0.39804    0.08565   4.647 5.18e-06 ***
---

```

```

Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

Residual standard error: 0.4756 on 281 degrees of freedom
Multiple R-squared:  0.09852,    Adjusted R-squared:  0.07285
F-statistic: 3.839 on 8 and 281 DF,  p-value: 0.0002641

```

The model selected through maximizing the $Z(\alpha=0.5)$ achieves a performance of 0.6897, which is higher than the full model's value of 0.6655. This improvement indicates that removing less informative predictors enhances forecasting accuracy. The optimal set of eight predictors—USR, EIF, ERECB, EFX, EDY, ESP, UIF, and URR—captures the majority of the predictive power while avoiding unnecessary noise from weaker variables. Statistically, the reduced model remains jointly significant (F-statistic $p < 0.001$), confirming that the selected predictors provide meaningful explanatory value. Several variables, such as ERECB, EDY, UIF, and URR, exhibit strong statistical significance. Although the adjusted R^2 is 7.3%, slightly lower than the full model.


```

# sub-question 2c: Search for optimal alpha that maximizes Z(alpha)
# Extract predictions from the best model found earlier
P_best <- predict(best$model, type = "response")

# Search grid for alpha
alpha_grid <- seq(0, 1, length.out = 2001)

Z_alpha_vec <- numeric(length(alpha_grid))

for (i in seq_along(alpha_grid)) {
  a <- alpha_grid[i]
  Z_t_a <- ifelse((P_best > a & returns_t > 0) |
                 (P_best <= a & returns_t <= 0), 1, 0)
  Z_alpha_vec[i] <- mean(Z_t_a)
}


# find maximizing alpha
idx_max <- which.max(Z_alpha_vec)
alpha_best <- alpha_grid[idx_max]
Z_best_alpha <- Z_alpha_vec[idx_max]

cat("Optimal alpha that maximizes Z(alpha):\n")
cat("alpha* =", alpha_best, "\n")
cat("Maximized Z(alpha*) =", Z_best_alpha, "\n")

# Plot Z(alpha)
plot(alpha_grid, Z_alpha_vec, type="l",
     main="Z(alpha) over alpha",
     xlab="alpha", ylab="Z(alpha)")
abline(v = alpha_best, col = "red", lwd = )

```

Optimal alpha that maximizes Z(alpha):
alpha* = 0.4945
Maximized Z(alpha*) = 0.6896552




3 Question 3 - CIR model for the term structure of interest rate

```
# load required libraries
suppressPackageStartupMessages({
  library(tidyverse)
  library(numDeriv)
  library(MASS)
})

data <- read_excel("data_coursework_Q3.xls") %>%
  rename(date = `...1`, r_US = `T r_$`, r_UK = `T r_£`) %>%
  arrange(date) %>%
  drop_na() %>%
  mutate(across(c(r_US, r_UK), ~ifelse(.x <= 0, NA_real_, .x))) %>%
  drop_na()

# plot interest rates of US and UK
library(ggplot2)
ggplot(data, aes(x = date)) +
  geom_line(aes(y = r_US, color = "US 1-month Rate")) +
  geom_line(aes(y = r_UK, color = "UK 1-month Rate")) +
  labs(title = "1-Month Interest Rates: US vs UK",
       x = "Date",
       y = "Interest Rate (%)") +
  scale_color_manual(values = c("US 1-month Rate" = "blue", "UK 1-month Rate" = "red")) +
  theme_minimal()
```



```
# Load required packages
suppressPackageStartupMessages(library(quantmod))
suppressPackageStartupMessages(library(MASS))
suppressPackageStartupMessages(library(Matrix))
suppressPackageStartupMessages(library(numDeriv))

# Extract interest rate data
r_US <- data$r_US # US 1-month rate
r_UK <- data$r_UK # UK 1-month rate

# Define log-likelihood for the CIR model
loglik_CIR <- function(theta, r) {
  mu <- theta[1]
  phi <- theta[2]
  sigma2 <- theta[3]

  T <- length(r)
  r_t <- r[2:T]
  r_t1 <- r[1:(T-1)]

  # Avoid numerical issues
  sigma2 <- max(sigma2, 1e-8)
  phi <- max(min(phi, 0.999), -0.999)

  # Calculate log-likelihood
  loglik <- -0.5 * log(2 * pi * r_t1 * sigma2) -
    (r_t - mu * (1 - phi) - phi * r_t1)^2 / (2 * r_t1 * sigma2)

  # Return negative log-likelihood (for minimization)
  return(-sum(loglik, na.rm = TRUE))
}
```

```

# Define log-likelihood for the Vasicek model
loglik_Vasicek <- function(theta, r) {
  mu <- theta[1]
  phi <- theta[2]
  sigma2 <- theta[3]

  T <- length(r)
  r_t <- r[2:T]
  r_t1 <- r[1:(T-1)]

  # Avoid numerical issues
  sigma2 <- max(sigma2, 1e-8)
  phi <- max(min(phi, 0.999), -0.999)

  # Calculate log-likelihood
  loglik <- -0.5 * log(2 * pi * sigma2) -
    (r_t - mu * (1 - phi) - phi * r_t1)^2 / (2 * sigma2)

  # Return negative log-likelihood
  return(-sum(loglik, na.rm = TRUE))
}

# Compute numerical Hessian
compute_hessian <- function(theta, loglik_func, r) {
  hess <- numDeriv::hessian(loglik_func, theta, r = r)
  return(hess)
}

# Estimation function
estimate_model <- function(r, model_name = "CIR") {
  # Initial parameter values
  init_params <- c(mean(r), 0.8, var(r))

  # Select model
  if (model_name == "CIR") {
    loglik_func <- loglik_CIR
    # Parameter constraints: mu > 0, 0 < phi < 1, sigma2 > 0
    lower_bounds <- c(1e-4, 1e-4, 1e-8)
    upper_bounds <- c(Inf, 0.999, Inf)
  } else {
    loglik_func <- loglik_Vasicek
    lower_bounds <- c(-Inf, -0.999, 1e-8)
    upper_bounds <- c(Inf, 0.999, Inf)
  }

  # Maximum likelihood estimation
  result <- optim(
    par = init_params,
    fn = loglik_func,
    r = r,
    method = "L-BFGS-B",
    lower = lower_bounds,
    upper = upper_bounds,
    hessian = TRUE
  )

  # Extract estimation results
  theta_hat <- result$par
  loglik_value <- -result$value

```

```

hessian <- result$hessian

# Calculate asymptotic covariance matrix
if (rcond(hessian) > 1e-12) {
  asymp_cov <- solve(hessian)
  # Ensure covariance matrix is symmetric positive definite
  asymp_cov <- nearPD(asymp_cov)$mat
} else {
  warning("Hessian matrix is near singular, using generalized inverse")
  asymp_cov <- ginv(hessian)
}

# Standard errors
std_errors <- sqrt(diag(asymp_cov))

# t-statistics
t_stats <- theta_hat / std_errors

# Results summary
results <- list(
  parameters = theta_hat,
  std_errors = std_errors,
  t_stats = t_stats,
  loglik = loglik_value,
  asymp_cov = asymp_cov,
  hessian = hessian,
  convergence = result$convergence
)

return(results)
}

# Model comparison function
compare_models <- function(cir_results, vasicek_results) {
  comparison <- data.frame(
    Model = c("CIR", "Vasicek"),
    LogLik = c(cir_results$loglik, vasicek_results$loglik),
    AIC = c(
      -2 * cir_results$loglik + 2 * 3,
      -2 * vasicek_results$loglik + 2 * 3
    ),
    BIC = c(
      -2 * cir_results$loglik + 3 * log(length(r_US) - 1),
      -2 * vasicek_results$loglik + 3 * log(length(r_US) - 1)
    )
  )

  return(comparison)
}

# Estimate models for US interest rates
cat("=== American Interest Rates ===\n")
cir_us <- estimate_model(r_US, "CIR")
vasicek_us <- estimate_model(r_US, "Vasicek")

# Output results
print_results <- function(results, country, model) {
  cat(sprintf("\n%s %s model estimation results:\n", country, model))

```

```

cat("Parameter estimates:\n")
estimates <- data.frame(
  Parameter = c("mu", "phi", "sigma2"),
  Estimate = results$parameters,
  StdError = results$std_errors,
  tStat = results$t_stats
)
print(estimates)
cat(sprintf("Log-likelihood: %.4f\n", results$loglik))
cat(sprintf("Convergence status: %d\n", results$convergence))
}

print_results(cir_us, "US", "CIR")
print_results(vasicek_us, "US", "Vasicek")

```

=== American Interest Rates ===

US CIR model estimation results:

Parameter estimates:

	Parameter	Estimate	StdError	tStat
1	mu	0.00010000	10.898336613	9.175712e-06
2	phi	0.99730042	0.002530594	3.940974e+02
3	sigma2	0.04373302	0.001089834	4.012816e+01

Log-likelihood: -236.9732

Convergence status: 0

US Vasicek model estimation results:

Parameter estimates:

	Parameter	Estimate	StdError	tStat
1	mu	5.8889468	2.92671969	2.012132
2	phi	0.9847641	0.01110412	88.684566
3	sigma2	0.5601000	0.04477204	12.510038

Log-likelihood: -353.4131

Convergence status: 0

```

cat("=== UK Interest Rates ===\n")
cir_uk <- estimate_model(r_uk, "CIR")
vasicek_uk <- estimate_model(r_uk, "Vasicek")

print_results(cir_uk, "UK", "CIR")
print_results(vasicek_uk, "UK", "Vasicek")

```

=== UK Interest Rates ===

UK CIR model estimation results:

Parameter estimates:

	Parameter	Estimate	StdError	tStat
1	mu	8.99627775	5.573881e-06	1.614006e+06
2	phi	0.99721507	6.987061e-03	1.427231e+02
3	sigma2	0.02468564	1.963883e-03	1.256981e+01

Log-likelihood: -198.1268

Convergence status: 0

UK Vasicek model estimation results:

Parameter estimates:

	Parameter	Estimate	StdError	tStat
1	mu	9.1633582	4.225711012	2.168477
2	phi	0.9929275	0.007994119	124.207245
3	sigma2	0.2585271	0.020669879	12.507433

Log-likelihood: -232.4193

Convergence status: 0

```
# Model comparison
cat("\n=== Model Comparison ===\n")
comparison_us <- compare_models(cir_us, vasicek_us)
comparison_uk <- compare_models(cir_uk, vasicek_uk)

cat("US Interest Rate Model Comparison:\n")
print(comparison_us)

cat("\nUK Interest Rate Model Comparison:\n")
print(comparison_uk)
```

=== Model Comparison ===

US Interest Rate Model Comparison:

	Model	LogLik	AIC	BIC
1	CIR	-236.9732	479.9464	491.1850
2	Vasicek	-353.4131	712.8262	724.0648

UK Interest Rate Model Comparison:

	Model	LogLik	AIC	BIC
1	CIR	-198.1268	402.2537	413.4923
2	Vasicek	-232.4193	470.8387	482.0773

```
# Output asymptotic covariance matrices
cat("\n=== Asymptotic covariance matrices ===\n")
cat("US CIR Model:\n")
print(cir_us$asymp_cov)

cat("\nUS Vasicek Model:\n")
print(vasicek_us$asymp_cov)

cat("\nUK CIR Model:\n")
print(cir_uk$asymp_cov)

cat("\nUK Vasicek Model:\n")
print(vasicek_uk$asymp_cov)
```

=== Asymptotic covariance matrices ===

US CIR Model:

3 x 3 Matrix of class "dpoMatrix"

	[,1]	[,2]	[,3]
[1,]	1.187737e+02	-2.533012e-02	8.096863e-04
[2,]	-2.533012e-02	6.403906e-06	-1.726767e-07
[3,]	8.096863e-04	-1.726767e-07	1.187737e-06

US Vasicek Model:

3 x 3 Matrix of class "dpoMatrix"

	[,1]	[,2]	[,3]
[1,]	8.565688e+00	-1.027894e-02	-2.998429e-06
[2,]	-1.027894e-02	1.233015e-04	4.344832e-09
[3,]	-2.998429e-06	4.344832e-09	2.004536e-03

UK CIR Model:

3 x 3 Matrix of class "dpoMatrix"

	[,1]	[,2]	[,3]
[1,]	3.106815e-11	-3.864257e-08	-2.060351e-11
[2,]	-3.864257e-08	4.881902e-05	4.451122e-09

```
[3,] -2.060351e-11  4.451122e-09  3.856837e-06
```

UK Vasicek Model:

3 x 3 Matrix of class "dpoMatrix"

```
      [,1]      [,2]      [,3]
[1,] 17.856633554 9.243223e-03 1.893588e-03
[2,]  0.009243223 6.390594e-05 9.795309e-07
[3,]  0.001893588 9.795309e-07 4.272439e-04
```

```
# Prediction function - Generate one-step-ahead predictions
generate_predictions <- function(model_results, r, model_type) {
  mu <- model_results$parameters[1]
  phi <- model_results$parameters[2]

  T <- length(r)
  r_t1 <- r[1:(T-1)] #

  if (model_type == "CIR") {
    # CIR model prediction: conditional expectation
    predictions <- mu * (1 - phi) + phi * r_t1
  } else {
    # Vasicek model prediction: conditional expectation
    predictions <- mu * (1 - phi) + phi * r_t1
  }

  return(predictions)
}

# Generate in-sample predictions
cir_pred_us <- generate_predictions(cir_us, r_US, "CIR")
vasicek_pred_us <- generate_predictions(vasicek_us, r_US, "Vasicek")

cir_pred_uk <- generate_predictions(cir_uk, r_UK, "CIR")
vasicek_pred_uk <- generate_predictions(vasicek_uk, r_UK, "Vasicek")

# Create prediction results data frame
actual_us <- r_US[2:length(r_US)]
actual_uk <- r_UK[2:length(r_UK)]

prediction_results <- data.frame(
  Date = data$date[2:nrow(data)],
  Actual_US = actual_us,
  CIR_Pred_US = cir_pred_us,
  Vasicek_Pred_US = vasicek_pred_us,
  Actual_UK = actual_uk,
  CIR_Pred_UK = cir_pred_uk,
  Vasicek_Pred_UK = vasicek_pred_uk
)

# plot predictions vs actuals for US
library(ggplot2)
ggplot(prediction_results, aes(x = Date)) +
  geom_line(aes(y = Actual_US, color = "Actual US Rate")) +
  geom_line(aes(y = CIR_Pred_US, color = "CIR Predicted US Rate")) +
  geom_line(aes(y = Vasicek_Pred_US, color = "Vasicek Predicted US Rate")) +
  labs(title = "US 1-Month Interest Rate: Actual vs Predicted",
       x = "Date",
       y = "Interest Rate (%)") +
  scale_color_manual(values = c("Actual US Rate" = "black",
```




```

    "CIR Predicted US Rate" = "blue",
    "Vasicek Predicted US Rate" = "red")) +

theme_minimal()

```




```

# plot prediction results for UK
ggplot(prediction_results, aes(x = Date)) +
  geom_line(aes(y = Actual_UK, color = "Actual UK Rate")) +
  geom_line(aes(y = CIR_Pred_UK, color = "CIR Predicted UK Rate")) +
  geom_line(aes(y = Vasicek_Pred_UK, color = "Vasicek Predicted UK Rate")) +
  labs(title = "UK 1-Month Interest Rate: Actual vs Predicted",
       x = "Date",
       y = "Interest Rate (%)") +
  scale_color_manual(values = c("Actual UK Rate" = "black",
                                "CIR Predicted UK Rate" = "blue",
                                "Vasicek Predicted UK Rate" = "red")) +

theme_minimal()

```



3.0.1 Analysis of CIR and Vasicek Models

Based on the estimation results, the CIR model demonstrates superior statistical fit compared to the Vasicek model in both the US and UK interest rate data, with higher log-likelihood values (US: -236.97 vs -353.41; UK: -198.13 vs -232.42) and lower AIC and BIC values, indicating that the CIR model can better capture the dynamics of interest rates, especially given its characteristic that volatility is related to the level of interest rates, which is more in line with the reality that interest rate fluctuations tend to be amplified at higher levels in the real financial market.

However, the parameter estimation of the US CIR model shows instability, such as the long-term mean being close to zero and having a very large standard error, which may reflect the significant fluctuations in US interest rates during the sample period since 1978 (such as during the oil crisis and periods of inflationary pressure), making model identification difficult.

In contrast, the parameter estimation of the UK CIR model is more stable, partly due to the relatively mild volatility of UK interest rates during this period, influenced by its stricter financial regulation and gradual monetary policy. Although the Vasicek model has more stable parameter estimation, its assumption of constant volatility limits its goodness of fit, especially in the highly volatile US environment.

Overall, the theoretical advantages of the CIR model are statistically verified, but its practical application requires careful handling of parameter uncertainty in light of the specific financial background of each country.